

Solving the Minimisation of the Open Stacks Problem

Joshua Davis

Monash University, Clayton,
Victoria, Australia, 3168

Supervisor: Maria Garcia de la Banda

School of Computer Science
and Software Engineering

Monash University, Clayton,
Victoria, Australia, 3168

2006

Contents

1	Introduction	1
1.1	The Problem	1
1.2	Previous Work	1
2	Methodology	2
2.1	Notation and Modelling	2
2.2	Preprocessing	2
2.3	Bounds	3
2.4	Search Algorithms	3
2.4.1	Exhaustive	3
2.4.2	Branch and Bound	4
2.4.3	A^*	5
2.4.4	Stepwise	7
2.4.5	Backwards	7
2.4.6	Binarychop	8
2.5	Non Opening Products	8
3	Implementation	9
3.1	Data types	9
3.1.1	Set Data Type	9
3.1.2	List Data Type	10
3.1.3	Lookup Table Data Type	10
3.2	Data Structures	10
3.3	Preprocessing	11
3.4	Search Algorithms	12
4	Guide to Using Software	13
5	Testing	14
6	Results	15

7	Discussion	19
7.1	Interpretation of Results	19
7.2	Optimisations	19
7.3	What I Learnt	20
7.4	Things I want to do	20
7.5	Did I reach my goals	20
8	References	21

1 Introduction

1.1 The Problem

A manufacturer makes a variety of products that are ordered by a number of customers in many combinations. However, the factory only has one machine which is capable of producing a single batch of a specific product at a time. Once the first product of a customer's order is produced, a stack is opened to which subsequent products in that order are stored. Following the completion of the last product, the stack is closed and products are shipped to the customer. Due to limited space in the factory, care must be taken to minimise the maximum number of simultaneous orders, hence to minimise the maximum number of open stacks.

Our ability to minimise the number of open stacks comes from finding an optimal schedule in which the products are to be produced. For example, if customers can choose from 10 products, it would yield $10!$ possible schedules. Our aim is to find the minimum number of open stacks for a given problem using the maximum of 100 products and 100 customers by employing a dynamic programming approach.

1.2 Previous Work

This project was largely based on our project leaders paper[1], where advanced algorithms implemented were discussed.

2 Methodology

2.1 Notation and Modelling

We will use the example depicted in table 1 to introduce the notation used throughout the report and to formalise our model. Let P be a set of products where $p \in P$ and C be a set of customers where $c \in C$. Let $c(p)$ be a function which returns a set of clients (column p from the table), ordering product p , we can extend this to where $P' \subseteq P$ we can now define $c(P')$ as

$$c(P') = \bigcup_{p \in P'} c(p)$$

We will need to evaluate a given schedule to implement simple search algorithms and to check that the solutions found by advanced (A*, stepwise, etc) search algorithms are correct. To help us do this will define an array, $schedule[i]$, which maps each point in time, $1 \leq i \leq |P|$, to an unique product p , where $p \in P$. Then it is possible to iterate over a schedule, however we need to evaluate how many stacks are active at time i to determine the minimum number of open stacks. Let $active(p, A)$ determine the number of active customers at each position i (i.e. $p = schedule[i]$) while retaining the maximum each possible time i . Let $A \subset P$ denote the products scheduled after $p = schedule[i]$, the set of active customers can now be defined as

$$active(p, A) = c(p) \cup \{c(P - A - \{p\}) \cap c(A)\}, \text{ where } p = schedule[i]$$

Schedule	1	2	3	4	5	6
Product	P_1	P_2	P_3	P_4	P_5	P_6
C_1	1	*	*	*	*	1
C_2	1	1	*	*	*	0
C_3	0	1	*	*	*	1
C_4	0	0	1	1	1	0
C_5	0	0	1	*	1	0

Table 1: The stars indicate already open stacks. At $schedule[3]$, P_3 alone has two active customers C_4 and C_5 as well as C_1 , C_2 and C_4 active yielding five open stacks. The current search space is $6!$ (720 possible schedule)

Now the problem may be formalised as,

$$\min_{each\ schedule} \max_{1 \leq i \leq |P|} \{active(p, A) | p = schedule[i]\}$$

2.2 Preprocessing

Several methods can be used to simplify the problem and thus reduce the amount of search space that needs exploring to find the optimal solution. Table 1. will be used to describe two methods, subsumption and partitioning which were implemented in the project. When these methods are successful it has been proven to dramatically reduce the time taken to find an optimal solution.

Subsumption eliminates scheduling products which have no impact on the optimal solution, iff $c(p') \subseteq c(p'')$. For example, if p' is a pen and p'' is a pen lid, customers ordering only p' , will also require p'' . Therefore p'' can be removed from the schedule and later reinserted directly after p' without affecting the optimal solution. Below is subsumed products removed from table 1, resulting in table 2.

Schedule	1	2	3	4
Product	P_1	P_2	P_3	P_6
C_1	1	*	*	1
C_2	1	1	*	0
C_3	0	1	*	1
C_4	0	0	1	0
C_5	0	0	1	0

Table 2: P_4 and P_5 have been removed, leaving a search space of $4!$ (24 possible schedules)

Partitioning the problem into sub problems, is possible if P can be partitioned into two sets P' and P'' , such that $c(P') \cap c(P'') = \emptyset$. Hence there are no customers requiring a product from P' and P'' . Then P' and P'' are independent problems and can be solved separately, where the solutions can be concatenated together. This dramatically reduces the search space from $|P|!$ to $|P'|! + |P''|!$, for example if $|P| = 10$, and $|P'| = 5$, the search space is reduced from millions to hundreds of possible schedules. Below table 3 has been partitioned.

Schedule	1	2	3	1
Product	P_1	P_2	P_6	P_3
C_1	1	*	1	0
C_2	1	1	0	0
C_3	0	1	1	0
C_4	0	0	0	1
C_5	0	0	0	1

Table 3: The double vertical lines indicates the partition in the problem, where the first partition has a search space of $3!$ (6 schedules) and the second has $1!$ totaling 6 possible schedules from 24 in table 2.

2.3 Bounds

The optimal solutions must exist within a domain, which can be trivially stated as the $\max_{p \in P} c(p)$ for the lower bound and $|C|$ for the upper bound. However the bound may be improved by various techniques such as examining the problem as a graph where the nodes are customers and edges products to find the maximum weighted clique, and taking that as the new upper bound. As the be tighten, part of the solution space may be avoided allowing the search to complete faster.

2.4 Search Algorithms

After preprocessing, the reduced problem is passed to a search algorithm where the optimal solution may be found. We will explore exhaustive, branch and bound, A^* tree search algorithms along with satisfiability variants stepwise, binarychop and backwards which expand on A^* . Ultimately, we will be able to apply knowledge about the stucture of the problem to reduce the amount of time searching.

2.4.1 Exhaustive

This algorithm exhaustively explores the entire search space, which is slow as every possible solution must be found and evaluated. This is not possible even with moderately sized $|P|$, as the computation required to find a solution explodes beyond the capabilities of current computers. However. it does have benefits over small values of $|P|$. Firstly, it reinforces the correctness of advanced algorithms as optimal solutions may be compared. Secondly, it provides a good metric to determine the efficiency of primitive methods by comparing differing implementations.

Algorithm 1 Exhaustive Search

```
procedure exhaustive( $S, \text{Sched}, U$ )  
  if  $S = \emptyset$  then  
    return maxOpen(Sched)  
  end if  
   $\text{min} := U$   
   $T := S$   
  while  $T \neq \emptyset$  do  
     $p := \text{any } p \in T$   
     $T := T - \{p\}$   
     $sp := \text{exhaustive}(S - \{p\}, \text{Sched} :: p, U)$   
    if  $sp < \text{min}$  then  
       $\text{min} := sp$   
    end if  
  end while  
  return  $\text{min}$   
end procedure
```

The algorithm can be thought of doing this; it begins by chopping the search space into smaller sub regions which the algorithm repeats recursively to form a tree. The algorithm is not concerned with narrowing the search space, therefore useful information is ignored at branches while the tree is grown, which will be utilized by later algorithms (branch and bound, etc) to speed up the search. The while loop selects a $p \in T$, which grows the tree horizontally, whereas the recursive function call, with p removed from S expands the tree vertically. When $S = \emptyset$, we are at a leaf of the tree with one possible $|P|!$ schedules. We evaluate the schedule, and the number of open stacks is passed up the tree being compared to other similarly obtained solutions at each leaves. This is done after the recursive call with the minimum (optimal solution) retained. Below is an example of a generated tree.




Figure 1: Exhaustive search tree of data in table 3

2.4.2 Branch and Bound

As $|P|$ increases the search space exponentially increases, so a system must be put in place to reduce the size of raw search performed. We will build on the idea of traversing a tree which was introduced in the above algorithm (branching). Unlike the exhaustive search algorithm, useful information contained at the branch is not simply ignored but is used to narrow the search space. This algorithm considers the information and decides whether the optimal solution resides further down that branch by comparing it to a bound which is continually updated as the tree is traversed (bounding). This avoids exploring superfluous search space which cannot improve on the current optimal solution (pruning).

Algorithm 2 Branch and Bound Search

procedure $bb(S, Sched, L, U)$ **if** $S = \emptyset$ **then** **return** 0 **end if** $min := U$ $T := S$ **while** $min > L$ **and** $T \neq \emptyset$ **do** $p := \text{index min}\{active(p, S - \{p\}) | p \in T\}$ $T := T - \{p\}$ **if** $active(p, S - \{p\}) \geq min$ **then** **break** **end if** $sp := \max(active(p, S - \{p\}), bb(S - \{p\}, Sched :: p, L, min))$ **if** $sp < min$ **then** $min := sp$ **end if** **end while** **return** min **end procedure**

The algorithm starts by checking if $S = \emptyset$, if so, no stacks are required and 0 is returned. Otherwise, it enters the while loop (branches) where a simple heuristic is used to determine the order of visiting each branch by picking $p \in T$ with the least active customers first. Hopefully this heuristic leads us to explore the optimal branch and subsequent iterations are avoided as the remaining p may have a greater or equal number of active customers compared to the current min (bounding). If this is true the while loop terminates as no better solution may be obtained at this branch (pruning). The maximum of the current active value compared to the value passing back up the tree is returned the optimal solution.




Figure 2: Branch and bound tree of data in table 3

2.4.3 A^*

Our current tree may have up to $|P|!$ branches explored in an attempt to find the optimal solution. After careful examination of the search tree it becomes evident that many branches throughout the tree contains pool of products where the next product to be selected is not unique (shown in figure 3). The active function used in the previous algorithm does not depend on the order of the products scheduled before and after the current branch. This implies the maximum number of open stacks at each of these repeated branches is the same. Therefore, the future does not depend on the past and now our problem is amenable to dynamic

programming. After evaluating a given branch the solution is stored and can be retrieved at branches with the same pool of products to be scheduled. This moves the search space down to $2^{|P|}$ as each of the unique branches make up the super set of all elements in P .




Figure 3: An example of a truncated tree with 4 products, only half the tree shown

The following is the algorithm.

Algorithm 3 A^* Search

procedure $stacks(S, stacks, L, U)$

if $S = \emptyset$ **then**

return 0

end if

if $stacks[S]$ **then**

return $stacks[S]$

end if

$min := U$

$T := S$

while $min > L$ **and** $T \neq \emptyset$ **do**

$p := \text{index min}\{\text{active}(p, S - \{p\}) \mid p \in T\}$

$T := T - \{p\}$

if $\text{active}(p, S - \{p\}) \geq min$ **then**

break

end if

$sp := \max(\text{active}(p, S - \{p\}), \text{bb}(S - \{p\}, \text{Sched} :: p, L, min))$

if $sp < min$ **then**

$min := sp$

end if

end while

$stacks[S] := min$

if $min > U$ **then**

$FAIL := FAIL \cup \{S\}$

end if

return min

end procedure

This algorithm is identical to branch and bound except solutions for each computed set are stored in a lookup table after the while, which can be later reused avoiding unnecessary recomputing by skipping the while loop.

2.4.4 Stepwise

The following three algorithms exploit that fact that A^* is very effective when the lower bound equals the upper bound. Stepwise, backwards and binarychop all turn our problem into a series of satisfiability problems by letting $U = L = n$ when calling *stacks*. If n is returned, n is the minimum number of stacks, otherwise *stacks* needs to be called again in hope to find the optimal solution. The following three algorithms step through the domain $[L, U]$ in different ways attempting to find the optimal solution. However, the lookup table after a non-successful call to *stacks*, the lookup table may contain viable solutions which can be reused in if stepping forwards. As we need to be able record what solutions are incorrect, A^* is slightly modified, after the while loop sets with values $> U$ are recorded, which are simply removed from the lookup table before the next call to *stacks*. The following is added after the while loop,

```

if  $min > U$  then
   $FAIL := FAIL \cup \{S\}$ 
end if

```

which is already included in algorithm 3 above.

Algorithm 4 Stepwise

```

procedure stepwise( $P, L, U$ )
  for  $try := L$  to  $U$  do
     $FAIL := \emptyset$ 
     $min := stacks(P, try, try)$ 
    if  $min = try$  then
      return  $min$ 
    end if
    for  $S \in FAIL$  do
       $stack[S] := 0$ 
    end for
  end for
  return  $U + 1$ 
end procedure

```

Stepwise, tries each value within the domain $L \leq try \leq U$ starting at the lower bound calling *stacks*(P, try, try) with *try*. Where *try* is incremented and failed sets removed from the table with each failure, until an optimal solution is found.

2.4.5 Backwards

Algorithm 5 Backwards

```

procedure backwards( $P, L, U$ )
   $gmin := U + 1$ 
   $try := U$ 
  while  $try \geq L$  do
     $min := stacks(P, try, try)$ 
    if  $min > try$  then
      return  $gmin$ 
    end if
     $gmin := min$ 
     $try := min - 1$ 
     $stack[S] := \emptyset$ 
  end while
end procedure

```

Backwards moves from U towards L , and has the ability to jump a number of values. A valid solution is found if we step one solution back too far and obtain an invalid solution. Unlike stepwise no solutions in the lookup table may be re-used.

2.4.6 Binarychop

Algorithm 6 Binarychop

```

procedure binarychop( $P, L, U$ )
   $gmin := U + 1$ 
  while  $L \leq U$  do
     $FAIL := \emptyset$ 
     $try := (U + L) \div 2$ 
     $min := stacks(P, try, try)$ 
    if  $min \leq try$  then
       $gmin := min$ 
       $U := min - 1$ 
       $stack[S] := \emptyset$ 
    else
       $L := try + 1$ 
      for  $S \in FAIL$  do
         $stack[S] := 0$ 
      end for
    end if
  end while
  return  $gmin$ 
end procedure

```

Binarychop is a variant of the two algorithms above, using both benefits of retaining solutions in the lookup table and the possibility of jumping a number of values. It is especially effective when the solution lays close to middle of the domain, as a binary search technique is used.

Unlike exhaustive and branch and bound at the completion of these algorithms we do not have a schedule. However, there is enough information in the lookup table to reproduce the schedule in polynomial time.

2.5 Non Opening Products

This is a technique used to reduce the search space in the previously explained algorithms (not including exhaustive), which returns a selected product before entering the while loop. It has been proven that products which do not open a new stack, if explored first does not affect the optimality of the solution. The proof is complicated and out of the scope of this project. This functionality can be added by the following lines before the while loop.

```

if  $\exists p \in S. c(p) \supseteq \{c(P - S) \cap c(S)\}$  then
   $alg(S - p, L, U)$ 
end if

```

3 Implementation

Here is a simple chart showing the program execution.

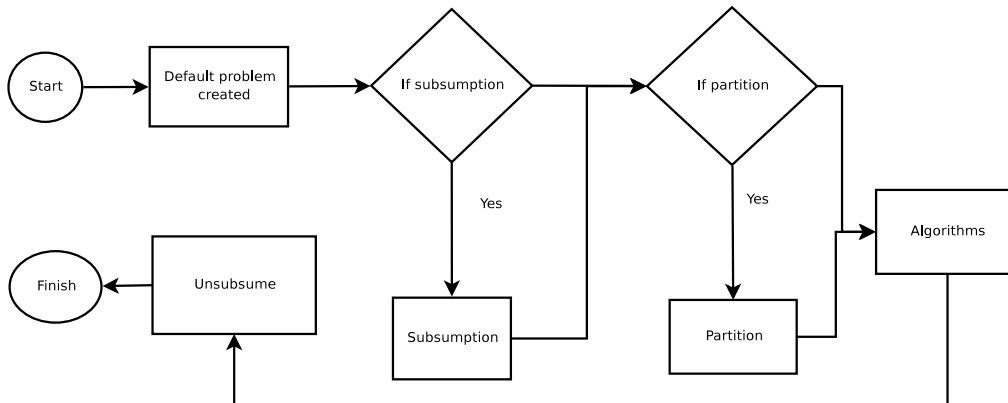


Figure 4: Program execution path

3.1 Data types

Implementing the above algorithms required many data types which may be done in many different ways, with varying strengths and weaknesses. Therefore, it was important that the most compatible implementation was matched with a suitable task to ensure maximum efficiency. If no thought was given to how the data type was used, the efficiency and perhaps effectiveness of obtaining a solution may have been lacking.

Initial analysis of the processes involved to obtain a solution showed that any operations performed during a search algorithm (staggering $O(|P|!)$) needs to be fast, but in operations performed in preprocessing (perhaps $O(n^3)$) is not as critical. This crucial observation lead to chosing implementations which favoured the operations preformed in the search algorithms. This lead to the following data types to be implemented in this project.

3.1.1 Set Data Type

The universal set is finite with up to 100 products and 100 clients, so using computer memory as bit string where the number of bits used is equal to the number of elements in the universal set (100 elements). A given set is then represented by a bit string in which corresponding elements in that set are 1 and all other bits are 0.

Set operations using a bitwise implementation are fast and simple in comparison to list operations, when the sets are of a small size as in this project. A desirable side effect of using bit strings is that bitwise operations $\&$ \mid \wedge can be used to implement the majority of operations. Furthermore, bitwise operations are fast, in most cases in constant time and complex procedures which may have been seen if using a list set representation are abolished. This may reduce bugs and lead to a more reliable code. Below is an example of an 8 bit segment of memory, where x can be 1 representing an element in the set or 0, the element is not in the set.

x x x x x x x Memory
7 6 5 4 3 2 1 0 Element

Here are two simple sets, in which the \mid operation computes the union.

1 0 0 1 0 0 0 1 = $\{0, 4, 7\} = A$
0 1 0 1 0 0 1 0 = $\{1, 4, 6\} = B$

$A \mid B \equiv A \cup B$
 $1\ 1\ 0\ 1\ 0\ 0\ 1\ 1 = 0,2,4,6,7$

3.1.2 List Data Type

Analysis of the search algorithms showed that the following list operations had to be fast, evident in while loop and functions calls from the body of exhaustive, and within the branch and bound algorithms. But the more advanced algorithms (A^* and backwards) lacked list operations all together, so less time was spent optimising the implementation. Choosing the list implementation proved difficult, below is a table comparing performance of differing implementations.

Operations	Array (end)	Array (Body)	Link List (end)	Link List (body)
Indexing	$O(1)$	$O(1)$	$O(1)$	$O(n)$
Inserting	$O(1)$	$O(n)$	$O(1)$	$O(n)$
Deleting	$O(1)$	$O(n)$	$O(1)$	$O(n)$

Table 4: Array vs. linked list implementations

Given that obtaining values via an index was more frequent than deleting and the majority of the operations were at the end of a list, the array implementation seemed like the best choice. Only limited analysis lead to this decision and analysis of the frequencies and locality of data lookups vs. list deletes was not explored. However, our chosen implementation only suffers slight performance loss using *maxOpen* which is used in the exhaustive and branch and bound algorithms. All the other algorithms used $O(1)$ operations.

3.1.3 Lookup Table Data Type

The obvious choice was to use a hash table with chaining. Not only are hash tables fast with all operations $O(1)$, given that collisions are avoided (or minimised). The hash table was used to mimic a linked list in recording failed sets in A^* . This is possible if the hash table is initialised with 1 bucket.

3.2 Data Structures

In order to successfully implement preprocessing and search algorithms all information required must be together in a structure. To further complicate matters, partitioning breaks the problem up into further sub problems, which alone has independent information. So the following structures were used.

```
typedef struct {
    char *n;          /* instance name */
    int cn;           /* customer no. */
    int pn;           /* product no. */
    int lb;           /* lower bound */
    int ub;           /* upper bound */
    int maxIterations;
    int *provenOptimal;
    int t;

    List *sp;         /* subsumed product list */
    Set **cust;       /* customer array N.B c(p) for fast look up */

    Segment *s;       /* contains all all MOSP's partitions and info */
    int sn;           /* contains the number of partitions */

    int flags;
}
```

```

Stats stats;

int **m;          /* matrix */
} Mosp;

```

In which all information is recorded. And,

```

/* contains everything an algorithm needs to complete */
typedef struct {
    int u;          /* contains the [u]pper bound */
    int l;          /* contains the [l]ower bound */
    int pn;         /* [p]roduct [n]umber */
    int cn;         /* [c]ustomer [n]umber */
    int *maxIterations;
    int *provenOptimal;
    int *t;

    List *p;        /* contains the [p]roduction */
    Set *up; /* contains all elements in production */
    Set *uc;
    Set **cust;     /* points to m->cust */
    Stats *stats;

    int flags;      /* launch time flags */
} Segment;

```

which contains information only relevant to sub-problems.

The following data structure was used to record all statistics.

3.3 Preprocessing

Reference to the code may help to to understand the following.

Determining what customers ordered a given product was frequent and expensive with an $O(|C|)$ method, so a lookup table was implemented. An array, $cust[i]$ returning a set of clients in constant time for product i .

Subsumption was implemented using two loops iterating over the product list. The outer loop initially starts with the first product in the list, cur where the inner loop always starts with the product at the front of current the list, cmp . The inner loop cycles to the last product before giving control back to the outer loop, where cur is moved the the next product in the list. This lets every product get compared with every other product. If the $cust[cur] \subseteq cust[cmp]$, cmp is deleted from the list and the loop continues until all subsumed products have been removed from the list.

The implementation of partitioning the problem into sub-problems is similar in that two loops are used to compare every product together. However, the outer loop is not a product, but a set of clients cmp , which is initialised to $cust[prodA]$ before the inner loop begins, also a new segment is made. If in the inner loop cmp is found to intersect with $cust[prodB]$, then $cmp = cmp \cup cust[prodB]$ and $prodB$ is then deleted from that list and added to the new segment and the inner loop starts again at the beginning of the list. The inner loop continues until no more products intersect, then cmp is set to the next product in the list and a new segment is made.

3.4 Search Algorithms

Reference to the code may help to to understand the following.

The exhaustive search algorithm described in methodology was efficient from an 1 to 1 implementation stand point. The only problems were extracting an optimal schedule and determining where variables should exist in the computer memory (stack vs heap). Where all sets have been allocated to the stack and lists to the heap (purely due to the way list were implemented, they need to go on the heap) . Using a structure to contain the optimal schedule and the minimum value appeared to be the most elegant solution, but was not used in favour of retaining the best schedule in a global variable (memory leak when using a structure to retain both prevented this method). At the base case the current *Sched* is placed in the global variable if it is better then the current best schedule in the global variable as well as the result from *maxOpen(sched)* being returned and passed back up the tree with the best retained.

The implementation of branch and bound was not clear like above, as the algorithm was not efficient with a lot of recomputing of active value within the while loop taking place. The sensible thing to do was to remove the active function call and to generate the *indexMin* data structure before the while loop. As *indexMin* contains the active value for each product, the multiple calls to active inside the while was no longer necessary. Retaining the optimal solution uses the same technique as in exhaustive search, but retaining the minimum was done in the while loop the same way as the algorithm in methodology.

Implementing A^* was particularly easy as all the code was the same as branch and bound expect for the hash function calls which needed adding before and after the while loops. Due to lack of time, I did not experient with special hash functions or hash tables sized with prime numbers.

Implementing stepwise, backwards and binary chop was also straight forward, with a direct translation from the algorithms started in methodology. Recording off successful sets was removed as it was redundant.

4 Guide to Using Software

The mosp utility runs both preprocessing and the search algorithm on a given problem instance and reports, statistics and the solution if obtained.

It has many command line arguments, based on a command line interface:

mosp flags problemFile

If no flags are specified or only -h is used, mosp will print flags which turn on different search algorithms, among other features.

```
-p find independent sub-problems
-d find redundant products, and consume them
-e solve using recursize exhaustive algorithm
-i <integer> stop after a number of iterations
-l <integer> use this number as the lower bound
-u <integer> use this number as the upper bound
-y solve using binarychop algorithm
-w solve using the backwards algorithm
-t show stats
-s solve using the stepwise algorithm
-f solve using initerative exhaustive algorithm
-b solving using branch and bound
-c use opening products first
```

The following example represents using subsumption and partitioning to solve a problem instance using the stepwise algorithm

mosp -dps problem

5 Testing

The testing was performed according to the following schemes. Depending on flags used during compilation, debugging, assertion and verbose output may be enabled. I am fairly confident that all the data type functions are correct.

All data type methods were given strict behaviour on all inputs, that is boundary cases were identified and dealt with elegantly (some instances the program terminates with an assertion). Unfortunately, in some complex functions the program terminates with an assertion due to speed considerations and shows the line number and source file where the program failed. This is not enough information in many cases to identify what caused the failure so extra information is included in the source code before the assertion describing why the assertion should be true.

All data type methods were extensively checked, with each boundary cases tested. The set library went under even more extensive tests, using randomly generated sets and testing relational, algebraic and special cases as this is the most used DT.

Due to lack of time, many higher order function lacked specific testing functions, but as these methods are built upon the data type methods, we can reason that is the logic is sound the behaviour of the function must be correct.

6 Results

The above implementation was tested with:

- Processor: Pentium IV 3.0Ghz
- Memory: 2.5GB
- Linux Kernel 2.6.9-34.ELsmp SMP
- Compiler: gcc 3.4.4
- Location: Monash University, student remote access terminal

The following tables are results obtained by running mosp with different search algorithms. Tables displaying full titles in each column were too big. Therefor shortened the following shortened titles stand for:

Instance, Instance Name. Optimal, Proven optimal. Parts, Partitions where each number is a partition and the number of products in it. Subs, Subsumptions. Open, Non-Opening Products counter. Fail, Fail set counter. Steps, Number of steps taken.

If all the values in a row are 0, the program could not complete due to exceeding 12mb. Only evident in \mathcal{A}^* .

Also the filenames were slightly changed to allow the tables to fit, GPN is GP_NWRS. War_1 is Warwick_1023. War_2 is Warwick_1718. War_3 is Warwick_1851. $Ware_1$ is Warwick_0003. $Ware_2$ is Warwick_0008.

Instance	Best	$ C $	$ P $	Optimal	Parts	Subs	Steps	Reused	Opening	Fail	Time(ms)
GP_1	0	50	50	0	42	8	1000001	0	0	0	293.143
GP_2	0	50	50	0	48	2	1000001	0	0	0	296.136
GP_3	0	50	50	0	50	0	1000001	0	0	0	298.412
GP_4	0	50	50	0	37	13	1000001	0	0	0	287.889
GP_5	0	100	100	0	100	0	1000001	0	0	0	297.535
GP_6	0	100	100	0	100	0	1000001	0	0	0	296.205
GP_7	0	100	100	0	99	1	1000001	0	0	0	310.19
GP_8	0	100	100	0	96	4	1000001	0	0	0	298.352
GPN_5	0	20	30	0	20	10	1000001	0	0	0	4122.9
GPN_6	0	20	30	0	23	7	1000001	0	0	0	5024.497
GPN_7	0	25	60	0	32	28	1000001	0	0	0	296.558
GPN_8	0	25	60	0	40	20	1000001	0	0	0	283.604
$Shaw_5$	0	20	20	0	18	2	1000001	0	0	0	3557.243
$Shaw_6$	0	20	20	0	19	1	1000001	0	0	0	3795.025
$Shaw_7$	0	20	20	0	18	2	1000001	0	0	0	3512.881
$Shaw_8$	0	20	20	0	19	1	1000001	0	0	0	3794.88
$Shaw_9$	0	20	20	0	18	2	1000001	0	0	0	3494.539
SP_1	0	25	25	0	16	9	1000001	0	0	0	2992.044
SP_2	0	50	50	0	44	6	1000001	0	0	0	309.396
War_1	0	20	20	0	16 1 1 1	1	1000001	0	0	0	2971.776
War_2	0	30	30	0	23 3 1 1 1	1	1000001	0	0	0	5028.988
War_3	0	30	30	0	12 1 3 1	13	1000001	0	0	0	2008.417
$Ware$	0	10	10	0	10	0	1000001	0	0	0	1610.626
$Ware_1$	3	10	10	1	7 3	0	13714	0	0	0	17.348
$Ware_2$	3	10	10	1	6 3	1	1971	0	0	0	4.333

Table 5: Exhaustive with partitioning and subsumption, total time taken 45s.231

Instance	Best	C	P	Optimal	Parts	Subs	Steps	Reused	Opening	Fail	Time(ms)
GP_1	45	50	50	1	42	8	408	0	0	0	21.727
GP_2	40	50	50	1	48	2	24867	0	0	0	1312.662
GP_3	40	50	50	1	50	0	5371	0	0	0	347.42
GP_4	30	50	50	1	37	13	94	0	0	0	5.975
GP_5	95	100	100	1	100	0	35196	0	0	0	8742.903
GP_6	0	100	100	0	100	0	1000001	0	0	0	115323.215
GP_7	0	100	100	0	99	1	1000001	0	0	0	185875.257
GP_8	0	100	100	0	96	4	1000001	0	0	0	93389.832
GPN_5	12	20	30	1	20	10	24799	0	0	0	198.407
GPN_6	12	20	30	1	23	7	253122	0	0	0	1817.737
GPN_7	0	25	60	0	32	28	1000001	0	0	0	13208.477
GPN_8	0	25	60	0	40	20	1000001	0	0	0	31298.675
$Shaw_5$	14	20	20	1	18	2	471516	0	0	0	3655.681
$Shaw_6$	14	20	20	1	19	1	31620	0	0	0	356.013
$Shaw_7$	13	20	20	1	18	2	29519	0	0	0	260.941
$Shaw_8$	14	20	20	1	19	1	64613	0	0	0	716.566
$Shaw_9$	13	20	20	1	18	2	13498	0	0	0	138.071
SP_1	9	25	25	1	16	9	769657	0	0	0	4294.533
SP_2	0	50	50	0	44	6	1000001	0	0	0	20514.9
War_1	3	20	20	1	16 1 1 1	1	64	0	0	0	2.474
War_2	4	30	30	1	23 3 1 1 1	1	31144	0	0	0	312.156
War_3	6	30	30	1	12 1 3 1	13	20	0	0	0	2.222
War_e	3	10	10	1	10	0	27	0	0	0	1.859
$Ware_1$	3	10	10	1	7 3	0	31	0	0	0	1.75
$Ware_2$	3	10	10	1	6 3	1	15	0	0	0	1.804

Table 6: Branch and bound with non-opening products, partitioning and subsumption, total time taken 8m1.829s

Instance	Best	C	P	Optimal	Parts	Subs	Steps	Reused	Opening	Fail	Time(ms)
GP_1	45	50	50	1	42	8	700	467	1982	0	70.618
GP_2	40	50	50	1	48	2	46821	43580	194366	0	2621.512
GP_3	40	50	50	1	50	0	38762	35929	148346	0	2097.014
GP_4	30	50	50	1	37	13	1656	1244	2886	0	90.657
GP_5	95	100	100	1	100	0	35485	33643	321356	0	5189.517
GP_6	75	100	100	1	100	0	407044	396710	4417382	0	65854.026
GP_7											80200.525
GP_8											86656.389
GPN_5	12	20	30	1	20	10	6124	4654	5169	0	157.643
GPN_6	12	20	30	1	23	7	15981	13593	15723	0	347.484
GPN_7	10	25	60	1	32	28	228697	183285	218004	0	4986.602
GPN_8	0	25	60	0	40	20	1000000	944706	1560918	0	25358.762
$Shaw_5$	14	20	20	1	18	2	7311	5720	4422	0	149.031
$Shaw_6$	14	20	20	1	19	1	10133	7981	5475	0	191.548
$Shaw_7$	13	20	20	1	18	2	5374	3901	3092	0	119.179
$Shaw_8$	14	20	20	1	19	1	12207	9950	7341	0	223.549
$Shaw_9$	13	20	20	1	18	2	4558	3316	2541	0	105.726
SP_1	9	25	25	1	16	9	12678	8231	822	0	166.169
SP_2	0	50	50	0	44	6	1000000	628304	272479	0	19229.413
War_1	3	20	20	1	16 1 1 1	1	621	245	132	0	74.989
War_2	4	30	30	1	23 3 1 1 1	1	12362	5341	2906	0	329.45
War_3	6	30	30	1	12 1 3 1	13	19	1	0	0	62.743
War_e	3	10	10	1	10	0	109	29	12	0	18.812
$Ware_1$	3	10	10	1	7 3	0	54	15	11	0	33.405
$Ware_2$	3	10	10	1	6 3	1	30	7	3	0	32.205

Table 7: $A\star$ with non-opening products, partitioning and subsumption, total time taken 4m54s

Instance	Best	C	P	Optimal	Parts	Subs	Steps	Reused	Opening	Fail	Time(ms)
GP_1	45	50	50	1	42	8	354	204	126	235	33.302
GP_2	40	50	50	1	48	2	3924	2826	1889	2941	152.844
GP_3	40	50	50	1	50	0	2545	1788	1725	2434	119.512
GP_4	30	50	50	1	37	13	64	20	37	44	21.886
GP_5	95	100	100	1	100	0	10500	8586	36816	38633	1041.119
GP_6	75	100	100	1	100	0	1268	679	789	1280	184.649
GP_7	75	100	100	1	99	1	4289	3209	4620	5602	358.734
GP_8	60	100	100	1	96	4	940	532	860	1173	135.873
GPN_5	12	20	30	1	20	10	891	535	188	527	35.906
GPN_6	12	20	30	1	23	7	1782	1250	644	1156	56.884
GPN_7	10	25	60	1	32	28	2623	1963	1264	1893	89.854
GPN_8	16	25	60	1	40	20	990740	888232	739103	841576	21928.027
$Shaw_5$	14	20	20	1	18	2	8681	6611	2766	4823	160.549
$Shaw_6$	14	20	20	1	19	1	5901	4095	915	2708	111.87
$Shaw_7$	13	20	20	1	18	2	3124	2107	781	1786	73.32
$Shaw_8$	14	20	20	1	19	1	7588	5647	1940	3868	141.236
$Shaw_9$	13	20	20	1	18	2	2807	1896	597	1495	66.308
SP_1	9	25	25	1	16	9	5200	3153	230	2262	97.219
SP_2	0	50	50	0	44	6	1000000	770448	93987	323539	19408.691
War_1	3	20	20	1	16 1 1 1	1	51	1	5	37	63.393
War_2	4	30	30	1	23 3 1 1 1	1	747	197	68	591	98.772
War_3	6	30	30	1	12 1 3 1	13	18	0	0	1	62.453
War_e	3	10	10	1	10	0	28	0	1	20	17.493
$Ware_1$	3	10	10	1	7 3	0	30	0	2	24	32.502
$Ware_2$	3	10	10	1	6 3	1	16	0	1	9	32.518

Table 8: Stepwise with non-opening products, partitioning and subsumption, total time taken 44.553

Instance	Best	C	P	Optimal	Parts	Subs	Steps	Reused	Opening	Fail	Time(ms)
GP_1	45	50	50	1	42	8	364	204	158	0	34.043
GP_2	40	50	50	1	48	2	4294	3313	5344	0	207.092
GP_3	40	50	50	1	50	0	3355	2511	3274	0	157.766
GP_4	30	50	50	1	37	13	150	40	89	0	26.398
GP_5	95	100	100	1	100	0	11648	10398	69075	0	1367.949
GP_6	75	100	100	1	100	0	78149	73969	532744	0	8485.849
GP_7	75	100	100	1	99	1	36519	32394	123564	0	2911.617
GP_8	60	100	100	1	96	4	209025	193862	672128	0	13327.962
GPN_5	12	20	30	1	20	10	529	317	131	0	28.435
GPN_6	12	20	30	1	23	7	1050	735	445	0	41.208
GPN_7	10	25	60	1	32	28	2644	1963	1275	0	82.711
GPN_8	16	25	60	1	40	20	627095	568016	540477	0	12098.695
$Shaw_5$	14	20	20	1	18	2	4442	3511	1821	0	93.047
$Shaw_6$	14	20	20	1	19	1	3367	2441	678	0	69.678
$Shaw_7$	13	20	20	1	18	2	1812	1283	553	0	50.738
$Shaw_8$	14	20	20	1	19	1	3917	3046	1222	0	79.231
$Shaw_9$	13	20	20	1	18	2	1555	1083	406	0	44.974
SP_1	9	25	25	1	16	9	5404	3199	231	0	93.001
SP_2	0	50	50	0	44	6	1000000	767845	172841	0	16402.846
War_1	3	20	20	1	16 1 1 1	1	106	4	21	0	65.836
War_2	4	30	30	1	23 3 1 1 1	1	1013	278	115	0	104.376
War_3	6	30	30	1	12 1 3 1	13	39	0	0	0	62.793
War_e	3	10	10	1	10	0	37	0	2	0	18.043
$Ware_1$	3	10	10	1	7 3	0	45	0	5	0	32.932
$Ware_2$	3	10	10	1	6 3	1	24	0	2	0	32.642

Table 9: Backwards with non-opening products, partitioning and subsumption, total time taken 55.948

Instance	Best	$ C $	$ P $	Optimal	Parts	Subs	Steps	Reused	Opening	Fail	Time(ms)
GP_1	45	50	50	1	42	8	364	204	158	235	22.5
GP_2	40	50	50	1	48	2	4542	3327	2950	4023	195.643
GP_3	40	50	50	1	50	0	3664	2673	3197	4040	177.85
GP_4	30	50	50	1	37	13	150	40	89	88	12.558
GP_5	95	100	100	1	100	0	12876	11062	69195	70711	1893.369
GP_6	75	100	100	1	100	0	13813	12562	74808	75660	1822.498
GP_7	75	100	100	1	99	1	14933	12365	28863	31036	1232.798
GP_8	60	100	100	1	96	4	22644	20199	62836	64802	1794.683
GPN_5	12	20	30	1	20	10	798	478	180	462	22.823
GPN_6	12	20	30	1	23	7	1559	1091	613	1014	44.007
GPN_7	10	25	60	1	32	28	2644	1963	1275	1893	93.349
GPN_8	16	25	60	1	40	20	867577	784261	712513	795711	24499.352
$Shaw_5$	14	20	20	1	18	2	7965	6176	2719	4475	166.714
$Shaw_6$	14	20	20	1	19	1	5122	3636	869	2320	101.298
$Shaw_7$	13	20	20	1	18	2	2714	1875	735	1541	74.334
$Shaw_8$	14	20	20	1	19	1	6336	4811	1749	3239	131.642
$Shaw_9$	13	20	20	1	18	2	1879	1287	446	986	42.318
SP_1	9	25	25	1	16	9	5361	3198	231	2347	107.482
SP_2	0	50	50	0	44	6	1000000	792030	192651	400571	18546.65
War_1	3	20	20	1	16 1 1 1	1	88	2	13	46	5.742
War_2	4	30	30	1	23 3 1 1 1	1	993	278	111	746	35.381
War_3	6	30	30	1	12 1 3 1	13	39	0	0	2	3.63
War_e	3	10	10	1	10	0	37	0	2	20	3.194
$Ware_1$	3	10	10	1	7 3	0	45	0	5	25	3.718
$Ware_2$	3	10	10	1	6 3	1	24	0	2	9	2.637

Table 10: Binarychop with non-opening products, partitioning and subsumption, total time taken 51.068

7 Discussion

7.1 Interpretation of Results

As expected the exhaustive algorithm failed to obtain an optimal solution for the majority of the tests, terminating early due to exceeding the 1 million step limit. This due to that fact that exhaustive search could only solve problems with < 9 products or the search space would exceed the 1 million allocated steps. To test this I choose 3 tests specifically with small search spaces $\leq 10!$, where two completed successfully, and the one equaling $10!$ failed.

Branch and bound, successfully obtained the solution to 19 of the 25 problems tested. This demonstrates that .bounding. and .pruning. is very effective at reducing the search space needing to be explored to obtain an optimal solution. It would have been interesting if I had time to implement an alternative heuristic for comparison. But it appears the choosing the product with the least active customers first was very successful, as it solved a further 17 problems on top of exhaustive with the same step limit. Also it appears the branch and bound is very efficient under some circumstances as it found the solution to GP_1 the fastest. I would assume GP_1 having a very obvious path to the solution and not requiring many past solutions needing re-use as it beat Astar and its variants.

Retaining partial solutions was the only enhancement added to branch and bound, perhaps the simplest change made to our program, yet allowed the ability to solve the majority of the selected problem instances (perhaps all of them but SP_2 if more than 12MB of memory could be used, and ≤ 1 million steps allowed). Table 7 shows 21 of the 25 tests obtained optimal solutions, which is a fair improvement upon the branch and bound. This really does go to show that simple solutions are always the best. It is interesting that if a solution was obtained, it was done well within the 1,000,000 steps allowed. This must be due to the nature of the search tree, if a solution was not obtained then the search must not be storing many used sets. Due to the limitation on memory usage two of the tests failed, which may have completed if allowed to finish executing.

Backwards, stepwise and binarychop all performed very well but it seems that how fast a solution is obtained is directly related to where the optimal solution lays in the domain from L to U. If close to L, stepwise is extremely fast, but also in some cases even when the best value is close to —P— stepwise still performs the best. This could be due to the fact that start at the back is harder than starting at the front retaining partial solutions as you step forwards. The opposite goes for backwards which appeared to be extremely fast when the best value was $\leq |P|/2$ indicated by table 9, tests GP_N_6 to SP_1 . Binarychop is extremely good at finding solutions if the lay in the middle, indicated by a single result GP_N_5 in table 10.

The fastest algorithm over all was stepwise, this is most likely due to its ability to continually reuse the stack lookup table. I guess the same logic can be applied from going to dynamic programming A^* from branch and bound, in that past solutions are constantly reused reducing recomputing.

7.2 Optimisations

After initial profiling both $O(n)$ methods in the set implementation, setCardinal and setNext accounted for up to 40% of the solution. So both these operations were optimised using lookup tables. Which improved the search time by a small margin.

The bit string sets also came in handy in quickly generate hash values, as the value of the string if read as an integer gives a high probability of being unique.

Retaining successful sets in A^* was ignored as backwards and binary when moving backwards, started with a new stacks table after each iteration.

The sliding window technique was used where ever possible, eg in *maxOpen*.

Avoiding excess calls to malloc and free was avoided by placing sets on the stack.

7.3 What I Learnt

What I learnt I have always struggled with recursion and knew I had to expose myself to it so that I could utilise it where it appeared more appropriate than iteration. Therefore, the natural thing to do was to choose MOSP as my 3rd year project as I would have no option but to learn how to use it correctly. Now that the project is complete and I have implemented four algorithms solely based on traversing a tree generated recursively. This project has allowed me to feel more comfortable with using recursion, however, I still believe I will require examples to work from when writing more complicated recursive algorithms. It is hoped with more time and practice I will find recursion as simple as iteration.

Keeping in mind the advice, 'use pen and paper before you touch the computer', I began my project on paper. Although, I would normally use the computer to find the best implementations for the datatypes, I found that using paper allowed me to be thorough and hence choose the best implementations. This was a valuable lesson and will be used as a starting point for further projects.

I reluctantly learnt to use long variable names, to which I have not grown a liking to you, as after a few weeks of not looking at the code, I quickly forgot what I was trying to do. Variable names which convey meaning facilitate the understanding of code.

I accidentally issued `rm *` in the source code directory 8 hours before the project was due, with my last backup 5 days old. So, I will always make frequent backups from now on.

7.4 Things I want to do

Unfortunately due to lack of time I was unable to implement a single heuristic of my own. Decreasing the upper bound would have been nice.

Also moving onto a 64 bit integer for the set DT would have been fun to see if it improved performance. Unfortunately I could not do this due to lack of time.

7.5 Did I reach my goals

I feel that I am happy with what I achieved, as always with a little more time, or perhaps better time management I could have. But all in all, I learnt a lot from this project, and I'm extremely lucky to have had Maria teach it to me. Thanks Maria for the semester!

8 References

- [1] M. Garcia de la Banda and P. J. Stuckey. Dynamic Programming to Minimize the Maximum Number of Open Stacks. IJCAI05 Constraint Modelling Challenge entry.